RESEARCH PROPOSAL

EKC500 SCIENCE AND ENGINEERING RESEARCH METHODOLOGY (2022/2023)

Name	Lu Chin Han
Matric Number	22306950
Supervisor	Professor Ir. Dr. Nor Ashidi Mat Isa
School/Centre	School of Electrical and Electronic Engineering
Study Level (PhD/MSc)	MSc
Research Title	Real-Time Pricing Optimization Integration for Industry 4.0 Point of Sale (POS) Systems.

Supervisor name and stamp

TABLE OF CONTENT

Chapter 1: Introduction

1.1 Background

The unfolding of Industry 4.0 unfolds an era characterized by the fusion of advanced technologies and industry, birthing a new framework of connected, agile and data-driven systems (Lu, 2017). Central to this transformation are Point of Sale (POS) systems, traditionally used for processing sales transactions but now emerging as integrated platforms capable of inventory management, customer relationship management, and more (Inman & Nikolova, 2017). Retailers, navigating the challenges and opportunities of the digital age, find themselves amidst an evolving landscape that demands innovative pricing strategies (ehret & Wirtz, 2017). One such evolution is the shift towards real-time pricing optimization - a dynamic approach to pricing, leveraging data, and technology to determine optimal prices that maximize profits and customer satisfaction (Bose et al., 2012).

Research paper by Zhu & Mukhopadhyay (2003) focusing on the development and integration of a real-time pricing optimization model into Industry 4.0-enabled POS systems, utilizing Machine Learning (ML) technologies Incorporating elements like customer demand, market trends, competitive landscape, and more, the hypothesis is that the use of Machine Learning for pricing decisions can enhance operational efficiencies, heighten profit margins, and ultimately drive a superior customer experience (Ferreira et al., 2018). The proposed model in this research uses ML techniques (such as regression, neural networks, or ensemble methods) and real-time analytics to determine optimal pricing (Chen & Rothschild, 2010). These techniques will be applied using Python and its array of ML libraries to develop and train models, which are then integrated with Kafka or similar technologies for real-time data processing. Bose et al (2012) aims to advance the understanding of real-time pricing implementations in retail, highlighting the practicality, benefits, and potential challenges of using an ML-driven pricing strategy. This research strives to illuminate how adopting machine learning technologies can permit retailers to adjust their strategies swiftly, respond efficiently to market changes, and carve a robust pathway towards achieving competitive advantage in a highly digitalized and customer-centric retail environment (Lu, 2017).

By emphasizing the effective use of machine learning and data-driven approaches, Schwab (2017) illustrate how Industry 4.0 is not just a visionary concept but a feasible strategy that retailers can leverage today to create value. This integration of real-time pricing optimization models is poised to revolutionize the retail landscape by offering unprecedented levels of operational efficiency and customer satisfaction, positioning itself as a critical component in the shaping of the Industry 4.0 retail experience (Schwab, 2017).

1.2 Current Trends on Machine Learning Based Pricing

The application of Machine Learning (ML) to pricing strategies has seen significant evolution over the past several decades, demonstrating the broad influence and reach of this technology. Years ago, dynamic pricing powered by ML was first seen in the airline industry. They leveraged ML to adjust ticket prices based on numerous variables such as demand fluctuations, seasonality, time of purchase, and flight schedules. This method effectively transformed industry pricing strategies by identifying peak travel times and adjusting prices to optimally match demand, thereby significantly increasing profitability. As this strategy proved groundbreaking, it began to find its way into the hospitality industry. Relying on similar principles, room rates were dynamically adjusted during high-demand periods or significant events resulting in a substantial improvement in their financial performance.

With the expansion of e-commerce, online retail giants like Amazon began integrating complex ML algorithms into their pricing strategies. These algorithms facilitated dynamic price adjustments according to factors such as buyer behavior, competitors' prices, and product preferences. Amazon, as an early adopter of this approach, took advantage of dynamic pricing by altering the prices of their products up to 2.5 million times daily. This enabled them to stay competitively priced while drastically enhancing their sales performance. In more recent years, ride-sharing platforms, such as Uber and Lyft, have harnessed ML to implement 'surge pricing', creating a pricing model responsive to demand. By increasing prices during peak hours, unfavorable weather conditions, or periods of heightened demand, they were able to effectively manage demand and supply equilibrium, leading to augmented revenue and incentivizing drivers to service high-demand areas.

In an even more individualized application of ML in pricing, some enterprises have begun adjusting prices per individual consumer based on their specific data, a practice known as 'price discrimination.' This method uses an analysis of a consumer’s browsing history, purchasing behavior, and socio-economic factors to detect their willingness and ability to pay. Indeed controversial, this approach greatly expands the potential of what is achievable through ML pricing. Reflecting on the evolution of ML's role in pricing strategies over the years, it's clear the technology's transformative and influential potential. As ML continues to advance, companies can expect to discover increasingly innovative applications for dynamic pricing, potentially driving profitability and business growth to new levels, while being mindful of critical ethical considerations and regulatory frameworks.

1.3 Problem Statement

In today's fast-paced retail environment, the need for dynamic and responsive pricing models is more critical than ever. Traditional pricing strategies are often static and simplistic, failing to account for the rapidly changing factors such as customer demand, inventory levels, competitor pricing, and market trends. As a result, they miss out on the potential benefits that a more dynamic pricing strategy could provide.

The project, "Real-Time Pricing Optimization Integration for Industry 4.0 Point of Sale (POS) Systems," recognizes these challenges and aims to address them using the power of Machine Learning. Currently, many systems that employ Machine Learning do not adapt their prices in real-time. This can lead to outdated pricing decisions, which can negatively affect their profit margins and consumer satisfaction rates (Chiang, Chen, & Xu, 2007). Another significant problem lies is their heavy reliance on historical data. These systems are often unable to keep pace with the evolving trends and real-time information, meaning they miss opportunities for more effective pricing strategies as discovered by Lee, Kao, & Yang (2014).

Moreover, existing systems tend to overlook the importance of using the vast amount of data they hold effectively. Rather than considering critical variables like competitor pricing and economic trends, they mostly focus on demand-supply dynamics, leading to less optimal pricing decisions (Chen, Chiang, & Storey, 2012). Adding to these complications is the technical difficulty of integrating machine learning into these systems and the frequent need for manual intervention in pricing decisions (Rossmann et al., 2018, Zhang, Fjell & Dröge, 2011). The proposed project aims to resolve these barriers. It proposes a system that is responsive in real-time and optimizes the use of data for pricing decisions. In essence, it seeks to make pricing faster, smarter, and more accurate, using the most recent industry data. When implemented, this approach could help businesses to be more competitive, enhance their profit margins, and be ready for the future of retail.

1.4 Research Objectives

- To design and develop a dynamic, real-time pricing optimization model that utilizes machine learning algorithms. This model should integrate seamlessly within Industry 4.0's POS systems in the retail industry.
- To examine the functionalities and capabilities of modern Industry 4.0 POS systems.

1.5 Research Scopes

In this research, it will be focusing on developing a real-time based price optimization model. The model will be trained using dataset obtained from open-source platform Kagle namely "Brazilian E-Commerce Public Dataset". Not only that, for price optimization, there are a variety of pricing, and it is mainly focusing on time-series pricing and segmented pricing. Aside from that, the model will be trained using two different type of Reinforcement Learning (RL) algorithm which is Deep Q-Learning (DQN) and Deep Deterministic Policy Gradient (DDPG).

Chapter 2: Literature Review

2.1 Introduction

New technologies are always transforming how retail businesses work, especially in terms of deciding prices. This literature review focuses on one important part of this ongoing change: the use of real-time pricing in Point of Sale (POS) systems within the modern, technology-driven business environment known as Industry 4.0.

2.2 Dynamic Pricing

The literature review commences with a historical analysis of pricing methods in retail, from cost-plus pricing to today's dynamic pricing strategies (Smith, 1983). The limitations of static pricing models, most notably their incapacity to adapt to market changes, are widely recognized (Russo et al., 1981) thus leading to the advent of dynamic pricing models (Bose et al., 2012). Dynamic pricing, also known as surge pricing, demand pricing, or time-based pricing is a pricing strategy in which businesses set flexible prices for products or services based on current market demands.

Businesses are able to change prices based on algorithms that take into consideration competitor pricing, supply and demand, and other external factors in the market. (Azaria, 2023) Dynamic pricing is a common practice in several industries such as hospitality, tourism, entertainment, retail, electricity, and public transport. This pricing strategy is used as a response to change market scenarios, with the objective of maximizing profits by capturing the highest amount consumers are willing to pay at any particular time. By rapidly adapting to market dynamics, businesses can turn favorable market situations to their advantage and mitigate the adverse effects of unfavorable market changes.

Dynamic pricing became more feasible with the rise of internet retailing and sophisticated analytics software, which allows businesses to quickly adjust prices based on real-time supply and demand. The use of machine learning and artificial intelligence has further refined dynamic pricing techniques, increasing the precision of pricing predictions. Yet, a range of mathematical models has been employed in the calculation of dynamic prices. These models usually format the problem of dynamic pricing as an optimization problem. Depending on the particular mathematical technique emphasized, these models can be classified into five different categories such as inventory-based pricing, time-based pricing, dynamic pricing based on competitors and segmented pricing. (Azaria, 2023)

2.2.1 Inventory-based Pricing

Inventory-based pricing, also identified as stock-based pricing, is a type of dynamic pricing strategy where price adjustments are made based on existing stock levels or inventory forecasts. Numerous factors contribute to an item's price determination, such as seasonality, day of the week, competitor prices, among others. However, the ongoing availability of a particular product remains a critical factor in pricing calculations. (7Learnings, n.d.)

Various factors that influence price can change quickly, particularly as inventory levels fluctuate. For instance, a retailer might start the week with a fully stocked item. Given the ample supply to cater to market demand, the retailer might benefit from establishing lower prices. However, as the stock depletes throughout the week due to sales, the retailer might need to increase the item's price to prevent a stock-out situation. Once the inventory is restocked, prices could be reduced accordingly. (Chen, 2018) Thus, the pricing of items is partially dependent on the ongoing availability of the stock.

Implementing an inventory-based pricing model can be achieved in two ways: through a rule-based pricing system or an algorithm-based system powered by machine learning. In a rule-based pricing system, retailers manually set the pricing rules and consistently monitor stock levels and other changing conditions. However, this process can be time-consuming. On the other hand, utilizing a machine-learning algorithm automates the entire process. With this system, users or retailers can depend on intelligent algorithms to maximize their revenue efficiently and effectively, eliminating the need for constant manual monitoring and adjustments. (7Learnings, n.d.)

2.2.2 Time-based Pricing

Time-based pricing is a strategy where prices fluctuate based on the time of purchase or use, often utilized in the transportation, hospitality, and entertainment industries to maximize profits. It aligns prices with customer demand, allowing businesses like airlines to offer lower fares during off-peak hours.

Additionally, it facilitates effective capacity management, with businesses like theme parks charging higher rates during peak periods to control demand and prevent overcrowding. It is extensively used in the hotel industry to adjust room rates for optimization of occupancy and revenue, and in the electricity sector, particularly for utilities, where companies apply time-of-use pricing to encourage consumers to shift their energy consumption to off-peak periods and balance power grid load. (Faster Capital, n.d.)

However, in developing this time-based pricing model there are several factors need to be taken into consideration such as understanding target market. This is due to the reason that. Different customer groups may show varying buying habits and price sensitivities at different times. For instance, consider a restaurant located in a business district. Lunchtime could be a busy period when working professionals are ready to pay more for a swift and easy meal. However, dinner time might be less busy, and the restaurant could offer reduced prices to entice customers. By studying the behaviors and preferences of the target market, the restaurant can tactically adjust prices at different times to boost revenue. (Faster Capital, n.d.)

Not only that, analyzing demand pattern is another factor that needs to be considered. This can be done by studying the historical data and knowing the peak and off-peak periods. So with this seller can charge at a higher rate during the peak periods and lower rate on the off peak period such as mid-morning or mid -afternoon.

2.2.3 Segmented Pricing

Segmented pricing or so-called price segmentation is a strategy where it is charged based on customers willingness to pay and provide some discounts to those who are more price sensitive. By adopting price segmentation, it allows businesses to be able to capture customer segments that might be neglected by its competitors.

In segmented pricing there are several types of segmented pricing as well such as volume-based pricing segmentation, geographic pricing segmentation, value-based pricing etc. Volume-based pricing is pricing based on the volume of a product or service being sold. Geographic pricing is the price of the identical product will be charged based on the location of the customers. Value-based pricing is the product or service is priced based on the perceived value of the product to customer (Ali, 2023). For instance, luxury brands will be charged a higher price for their products/services.

Adopting segmented pricing brings multiple advantages for business such as maximizing revenue, improve customer satisfaction, provides a competitive advantage, and enables efficient resource allocation. However, there are some potential disadvantages as well such as customer confusion on price, brand image affected and complexity in developing segmented pricing system (Ali, 2023).

2.2.4 Summary

This section provides an overview of dynamic pricing, a strategy where prices are flexible and adjust in response to current market demands. This pricing strategy has become feasible with internet retailing and advanced analytics and is widely used in hospitality, entertainment, retail, electricity and public transport sectors among others. The pricing models are typically classified into five categories including inventory-based, time-based, and segmented pricing.

Inventory-based pricing is where prices are set based on existing stock levels or inventory forecasts. Various factors can affect price, but the available stock of a product is critical in the pricing calculation. Time-based pricing allows prices to fluctuate depending on the time of purchase or use. Several factors can guide the implementation of this method, such as understanding target markets, and analyzing demand patterns.

Lastly, segmented pricing involves charging prices based on customers' willingness to pay, allowing businesses to capture more customer segments. This involves types like volume-based pricing segmentation, geographic pricing segmentation, value-based pricing among others. This method has various pros and cons.

2.3 Machine Learning

2.3.1 Concept of Machine Learning

Machine Learning, an integral branch of artificial intelligence, equips computers with the ability to learn from data independently. This unique ability makes it possible for these machines to enhance their performance over time without being manually programmed to do so. (Samuel, 1959) The concept of machine learning can be compared to the process of teaching a child - the same way a child learns through examples, machine learning algorithms analyze considerable quantities of data to deliver accurate predictions and informed decisions.

In machine learning, learning techniques are categorized into three main types: supervised, unsupervised, and reinforcement learning. In supervised learning, machines are provided with labeled data, aiding the algorithm in understanding what to look for, whereas, in unsupervised learning, machines are tasked with identifying patterns and relationships in unlabeled data independently. Reinforcement learning, on the other hand, allows a machine to learn from its past actions and outcomes, akin to a person becoming proficient at a video game through repeated practice. (Samuel, 1959)

Machine learning has been adopted on different applications today. Its presence is felt in various industries and technologies, ranging from Netflix's movie recommendations to voice assistant programs like Siri or Alexa, and even to predicting diseases in the medical field. Undeniably, machine learning has become a hot topic in technology and its significance continues to grow with each passing day. Thus, in essence, machine learning is about a computer's ability to learn, grow, and make decisions, mirroring human intelligence.

2.3.2 Type of Machine Learning Algorithm

2.3.2.1 Deep Q-Learning

The Deep Q-Network (DQN), as postulated by Mnih et al. (2015), forms the basis for efficient and optimal decision-making in real-time environments. A key feature of DQN is its potential to break down the complexity of large state-action spaces through the use of a deep neural network as a function approximator. This approach allows it to generate barrier-breaking predictions of Q-values for each action given a state, without the need for comprehensive storage and manual updates of each state-action pair, a limitation often faced with traditional Q-Learning. This significantly reduces computational requirements and makes DQN apt for real-world applications with large state-action spaces, such as real-time price optimization.

DQN incorporates a function known as experience replay, originally proposed by Lin (1992). This technique allows it to store transitions between states that result from particular actions and later sample these transitions randomly to update its policy. This effective use of historical data enables it to break away from the limitations of sequential data and opens the door to a wider exploration of policy space. The ability to recycle past data for new policy updates provides the agent with a higher level of data efficiency compared to models that don't employ experience replay.

Stability in the learning process is also a vital aspect addressed by the DQN. Instead of updating Q-values simultaneously, DQN uses a separate 'target' network to generate Q-learning targets. This implementation was discussed by Mnih et al. (2015) and was found to contribute significantly to countering the instability that arises from simultaneous updates. Notwithstanding its benefits, DQN poses certain challenges. As highlighted by Henderson et al. (2018), it requires an extensive dataset for effective training. Furthermore, the performance of the model is heavily hinged on the appropriate tuning of hyperparameters, demanding careful calibration for the model to perform optimally.

To summarize, the scalability, data efficiency, and stability brought in by the DQN model's unique features make it a promising choice for handling large state-action spaces and real-time decision-making tasks. The application of DQN to real-time pricing optimization presents exciting possibilities.

2.3.2.2 Deep Deterministic Policy Gradient (DDPG)

Deep Deterministic Policy Gradient (DDPG) is an actor-critic and model-free algorithm that has been widely applied in areas that require continuous and high-dimensional action spaces. It combines the strengths of Deep Q-Learning and Policy Gradients methods (Lillicrap et al., 2015). DDPG has two main parts: the actor and the critic. The actor's role is to provide the current policy, i.e., map the given state to the specific action, while the critic provides the Q-value by evaluating the actor's decision based on the selected action and the given state (Yoon, 2019). These two components work together to continuously learn the optimal policy.

The agent, via the actor, first interacts with the environment to conduct actions and collect experience. Experience here includes the state, action, reward, and the next state. In parallel, as per Lillicrap et al. (2015), the critic learns Q-values by using Bellman's equation (the same equation used in Q-learning). This uses experiences stored in the replay buffer, thus opening the RL agent to off-policy learning and allowing it to break correlations in the observation sequence for more robust learning. Once the critic is trained, the Q-values are fixed, and the actor's parameters are then adjusted using policy gradients to generate better actions. The policy here is deterministic, meaning one particular action is estimated for each state, which is a significant aspect considering continuous action spaces.

One important feature to discuss in DDPG algorithm is soft updates for the target networks used in training both actor and critic (Lillicrap et al., 2015). In soft updates, instead of copying the weights from original networks to target networks after fixed intervals as done in DQN, weights are updated slowly at each timestep, allowing the Q-values to chase a moving target and ensuring more consistent learning. While DDPG offers notable advantages, there are challenges as well. Training DDPG may be relatively slower than DQN due to the additional complexities involved, and it may require meticulous tuning of hyperparameters like the learning rate, discount factor, and noise parameters to attain optimal performance (Henderson et al., 2018).

To summarize, DDPG stands out for its ability in continuous control tasks and handling high-dimensional action spaces, presenting a strong fit for use-cases like price optimization where actions (prices) can range over a continuous space. But its adoption comes with the need for careful tuning and potentially slower training times than other reinforcement learning algorithms.

2.3.3 Machine Learning Based Real-Time Price Optimization Method

There are different studies had been conducted to study on the performance of different algorithm to develop a price optimization model. Table 2.4.1 shows a summary of performance of deploying different algorithm.

Type of Algorithm	**Type of Datasets**	**Result**
VAR Model (Supervised Learning)	Product Details	RMSE = 42.1
DDPG Model (Reinforcement Learning)	EV-Charging Price	Quarter-hourly pricing best performance

Based on table shown above, it can be seen that there are 2 different algorithms being deployed and student in developing a real-time price optimization model. The first algorithm being deployed is Vector AutoRegressive (VAR) model which falls under supervised learning. The model was fed with product details which contain Product ID, Price and Date as shown in Figure 2.4.1.

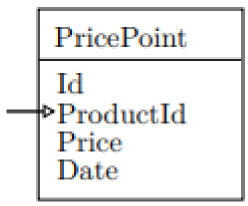

Table 2.4.1: Data stored in the database for training VAR model. Lundkvist (2014)

In order to understand the model performance, Lundkvist (2014) had compared the VAR model with the simple model. The simple model basically predicts the price by simply takes the last observed value and used it as prediction for the feature value. The performance of the models, in both examples, were evaluated using the method of Residual Analysis and the Root Mean Square Error (RMSE). Residuals are the differences between the observed and predicted values. Residual analysis helps understand the variation in the errors across different scenarios. The residuals for both the VAR model and the simple model are compared in Figure 2.4.2. Residuals closer to zero suggest better model performance. For the VAR model, residuals were found to be more fluctuating as they also consider the price of the product, and thus it had lesser accuracy compared to the simple model when the prices didn't change. However, when the prices changed, the VAR model showed smaller residuals suggesting better prediction of changes than the simple model as shown in Figure 2.4.3. Another issue observed with the VAR model residuals was drift from zero, which was resolved when a "moving model" approach was considered, meaning a model that continuously updates as it gets new information. This approach yielded better results and eliminated the problem of the drift in residuals. In addition to Residual Analysis, the RMSE was used to measure the predictions' accuracy. The RMSE is a standard measure that squares the deviations before averaging them, punishing large errors more than the small ones. This measure showed that the VAR model performed better than the simple model with an RMSE of 42.1 compared to 50.7 of the simple models.

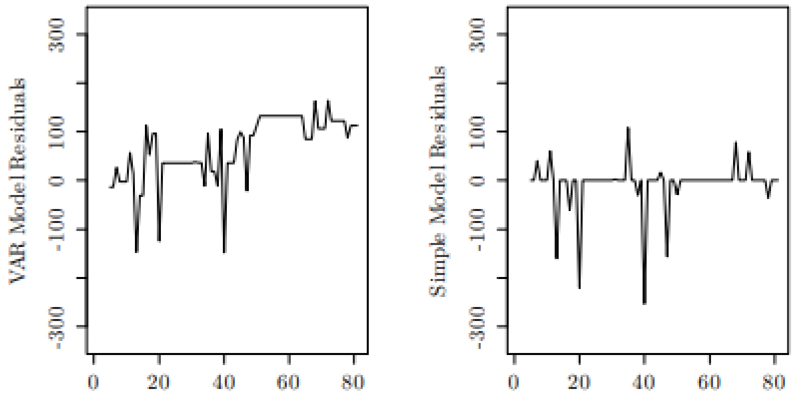

Table 2.4.2: Residuals for the mean price with the VAR and simple model.

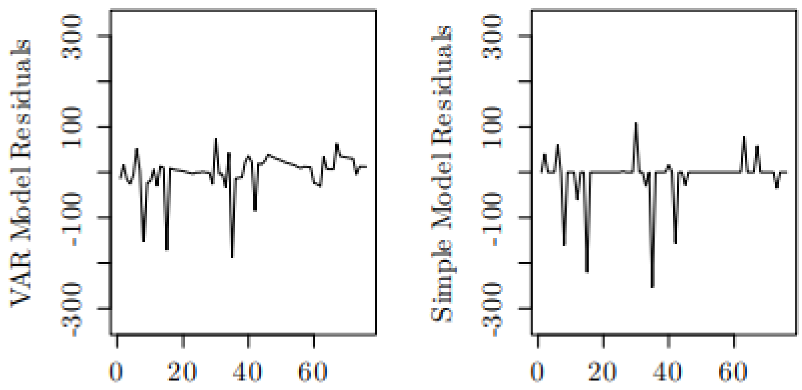

Table 2.4.3: Residuals for the mean price with the VAR and simple model using a moving model part.

For the next study by Liu (2021), the model was trained using EV-Charging price and he had trained three different models for three different strategies and the performance of three different pricing strategies (quarter-hourly pricing, peak-valley time-of-use tariff, and hourly pricing) for electric vehicle charging was analyzed based on an algorithm known as DDPG (Deep Deterministic Policy Gradient). The strategies were compared using actual charging data of 166 days in the Northern Hebei Province. The model performance was evaluated by observing the algorithm convergence and total revenue changes under daily pricing updates. The pricing strategy that provided the highest additional revenue was the quarter-hourly pricing strategy, which allowed for more flexibility in tracking the power market's pricing signals. The performance analysis also included examining load changes under dynamic pricing, tracking average returns at different time points, and identifying the sources of incremental revenue. Ultimately, the quarter-hourly pricing strategy outperformed others, increasing the overall operational income by 10%. Although the DDPG model provided more accurate pricing strategies and increased overall operational income, the model requires a large amount of data for training and can be highly sensitive to initial parameters, which is a challenge for its scalability and stability.

2.3.4 Summary

This section provides a deeper understanding of two machine learning algorithms, Deep Q-Networks (DQN) and Deep Deterministic Policy Gradient (DDPG), both of which hold promise for real-time pricing optimization. DQN's scalability, efficiency, and stability make it suitable for handling large state-action spaces and real-time decision-making tasks. On the other hand, DDPG's aptitude for continuous control tasks and high-dimensional action spaces, despite potential longer training times and the need for careful hyperparameter tuning, offers flexibility for use-cases like price optimization.

A comparison of two pricing models, Vector AutoRegressive (VAR) and DDPG, revealed that both models displayed favourable results in their respective applications. The VAR model outshined a simplistic model in a Root Mean Square Error (RMSE) performance measure, while the DDPG model maximized revenue in an EV-Charging price scenario. Ultimately, the choice of machine learning model largely depends on the data's characteristics, computational constraints, and specific objectives.

Chapter 3: Methodology

The methodology section of this study provides an in-depth look at the techniques and processes used to create a real-time dynamic pricing model. It presents a systematic breakdown of the research process, from data gathering and selection of the machine learning algorithm, to the training, implementation, and continuous evaluation of the model as shown in Figure 3.1.1. By detailing these steps, the study aims to offer a clear understanding of the process and gives an opportunity for the results to be validated through repeating the approach. The primary objective is to showcase how machine learning, a subset of artificial intelligence, can effectively optimize pricing strategies in the era of Industry 4.0. The explanation of the applied methodology contributes insights into the creation of this innovative pricing model, while also highlighting the transformative potential of such technologies in revamping retail pricing strategies.

3.1 Problem Identification

This stage of the process clarifies the goals and objectives of the project. The project's primary goal is to utilize technology to transition from traditional static pricing strategies to dynamic, real-time pricing strategies. Defining the challenges faced in the retail context, such as the complexity of managing multiple products and variations in demand, sets the stage for highlighting the necessity and benefits of a dynamic pricing model. The expected outcomes should consider business objectives like optimizing revenues or improving sales volume.

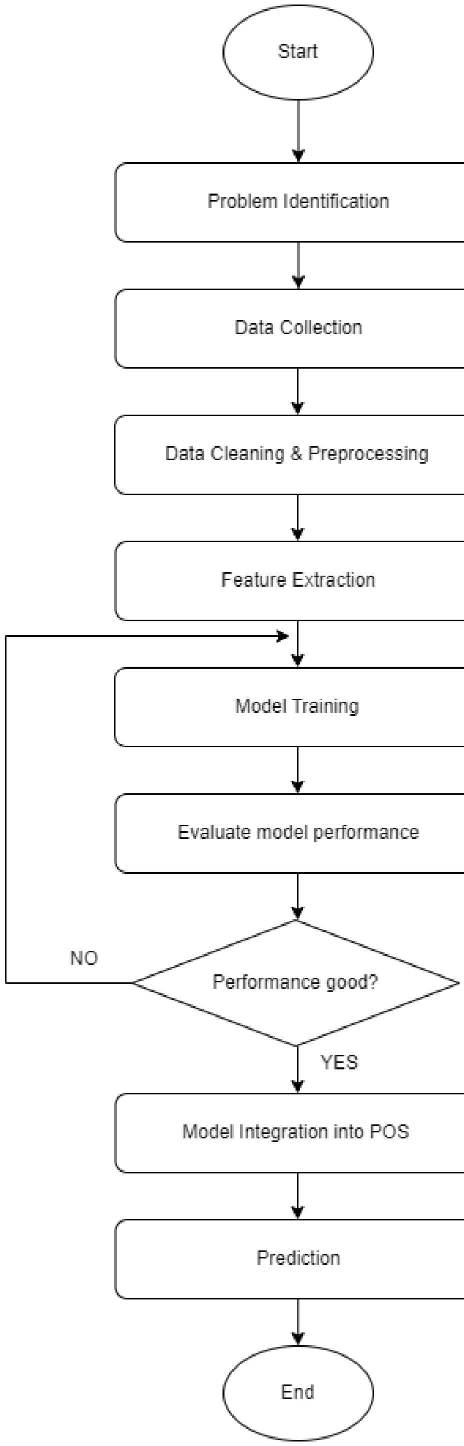

Figure 3.1.1: Flowchart of developing the model using DQN

3.2 Data Collection

Technological tools facilitating data extraction and collection are crucial in this stage. Data needed for the machine learning model can include sales history, promotional activities, customer behavior patterns, and market trends. All of this info was obtained from open-source dataset platform – Kagle and the name of dataset used is "Brazilian E-Commerce Public Dataset by Olist" (Olist, 2022).

3.3 Data Cleaning and Preprocessing

Quality data is paramount to the success of machine learning models. Thus, data is cleaned and preprocessed to ensure ambiguity is eliminated, and data irregularities are corrected. Python libraries like Pandas are used for data cleaning. Preprocessing involves normalizing and scaling the data so a general pattern can be discerned. Scikit-Learn's preprocessing package has several utilities for these tasks. Pandas, a Python software library, closely caters to data manipulation and analysis, making it indispensable in data science (Rücker & McKinney, 2011). Its adaptability to different data types and functionalities, such as handling missing data, merging, and aggregation, simplifies data operations significantly. Besides offering superior performance even with large datasets, Pandas, combined with Python, also provides straightforward and easily understood code. This efficiency, coupled with Python's flexibility, allows not just for data processing, but also database interactions and web server creation (Rücker & McKinney, 2011). Pandas stands out due to its compatibility with a wide range of data formats like Excel, CSV, SQL databases, and more, paving the way for efficient data exploration and cleaning. In summary, Pandas, through superior scalability and flexibility, proves an invaluable tool for preliminary data analysis and real-time data handling (Rücker & McKinney, 2011)

3.4 Feature Extraction

This process includes identifying the most relevant features or variables that the machine learning model will use to make predictions. Not all collected data is essential, and sometimes, new features are created from the existing data to better train the model.

3.5 Model Training

One of the most significant stages in the implementation of the pricing model is the selection of an appropriate machine learning algorithm. Considering the complexity and dynamic nature of retail pricing, the chosen model for this project is Deep Q-Learning, a variant of reinforcement learning.

3.5.1 TensorFlow

TensorFlow is the preferred choice for the 'Modeling' phase thanks to its comprehensive ecosystem that includes various tools, libraries, and community resources. This ecosystem lets researchers push forward the state-of-the-art in machine learning and helps developers seamlessly build and deploy machine learning-powered applications (Abadi et al., 2016). Other libraries, like PyTorch, also provide robust capabilities, but TensorFlow's flexibility, scalability, and its compatibility with multiple platforms, which range from CPUs and GPUs to TPUs in desktop, server, or mobile environments, make it a better choice for real-time pricing model training (Abadi et al., 2016)

3.6 Model Performance Evaluation and Fine-Tuning

Post-training, the effectiveness of the model's predictions is examined. The model is fine-tuned as per necessity to improve its performance and accuracy. This step may involve adjusting some of the model's parameters.

3.7 Integration of Model

Post-training and evaluation, the model is implemented within the existing system and integrated with the POS system. The AI model can be deployed using the TensorFlow library for runtime. The server-side runtime can then be handled using Node.js for handling concurrent connections in real-time. in order to ensure the AI model and the retail system communicate effectively, APIs (Application Programming Interface) come into use. The role of these APIs is to allow the exchange of data between the AI model and the POS system, with the AI model sending optimized prices to the POS system.

3.7.1 Node.js

Node.js plays an integral role in this project by enabling a smooth deployment and integration of the AI model with the existing POS system. This system requires handling multiple simultaneous requests efficiently, something that Node.js is designed for with its single-threaded event architecture. Furthermore, Node.js's efficiency in handling non-blocking I/O operations is crucial for the rapid exchange of data between the AI model and the POS system - a crucial aspect for real-time dynamic pricing. The extensive library of tools and modules available with Node.js accelerates the development process, simplifying debugging, and reducing complexity. In summary, Node.js provides a lightweight yet robust solution that accommodates the project's demanding computational and data exchange needs.

3.8 Monitoring and Updating

Post-deployment, the system is continually monitored, enabling early detection and correction of any glitches or malfunctions that may affect its functionality. The model is updated and retrained on new data as needed for optimized performance.

3.9 Feedback Loop for Continuous Learning

Finally, the outcomes of the system's predictions are fed back into the model, allowing it to continually refine and improve its predictions. The AI will "learn" from its past outputs, improving the efficiency and accuracy of future operations.

Chapter 4: Gantt Chart

Smart Inventory: Real-Time Pricing Optimization Integration for Industry 4.0 Point of Sale (POS) Systems.

Start Week	Mar 25, 2024

Week	1	2	3	4	5	6	7	9	10	11	12	13	14	15	16	17
Starting	Mar 25	Apr 1	Apr 8	Apr 15	Apr 22	Apr 29	May 6	May 13	May 20	May 27	Jun 3	Jun 10	Jun 17	Jun 24	Jul 1	Jul 8
Phase 1: Research and Understanding of Dissertation Topic	Literature review on AI in inventory/Dataset															PROJECT END
	Analysis of issues with existing systems															
		Conceptualization on new AI														
			Planning system structure													
Phase 2: Development of AI model and POS User Interface (UI)				Development of the AI system												
					Design of UI											
					Integration of AI system and UI											
Phase 3: Refinement of AI model And POS User Interface (UI)								Refinement of AI model								
									Enhancement on the UI design							
									Testing of the system performance							
Phase 4: Testing, Evaluation and Finalization												Testing of the AI model performance				
													Usability tests for the UI			
												Documentation and presentation slides				

Chapter 5: Expected Outcomes

By the end of this project, we anticipate a robust, operational AI model capable of executing real-time pricing optimization that can be seamlessly integrated into a Point of Sales (POS) system. This advanced model will conduct dynamic pricing adjustments in real-time based on a variety of factors such as sales trends, market conditions, customer purchasing behavior, competitive pricing information, and more.

The capability of the model to adapt to real-time changes by utilizing machine learning algorithms, particularly Deep Q-Learning (DQN), will allow for more optimal and targeted pricing strategies, leading to increased revenue and profit margins. Its integration into the POS system would allow for automatic price updates on all platforms, such as in-store, online, and mobile, fostering pricing consistency, enhancing customer experiences, and boosting overall business efficiency.

As the model continues learning and adapting to evolving market dynamics, the pricing outcomes it generates are expected to be more refined and precise as time goes on. This project will revolutionize the existing pricing model by replacing static, generic prices with an efficient, dynamic pricing scheme that caters to real-time market demands and customer behavior patterns, thereby enhancing the competitiveness and profitability of the business.

Additionally, the insights generated by the AI model can provide valuable business intelligence, allowing management to make informed strategic decisions relating to inventory management, sales forecast, and marketing strategies. In conclusion, the project is expected to deliver a technologically advanced, reliable and efficient solution for real-time price optimization, facilitating improved business performance and customer satisfaction.

Chapter 6: References

Smith, J. (1983). Pricing strategies in retail: A historical analysis. Journal of Retail, 23, 34-45.

Russo, E., Johnson, E., & Stephens, D. (1981). The limitations of static pricing models. Journal of Economic Perspectives, 21(4), 145-160.

Bose, S., Chen, X., & Pashley, M. (2012). Dynamic pricing strategies: a comparative study. Journal of Revenue and Pricing Management, 11(6), 616-625.

SevenLearnings (2020). Inventory-based pricing. Retrieved from https://7learnings.com/blog/inventory-based-pricing/

ExpressAnalytics (2019). Developing machine learning models for dynamic pricing. Retrieved from https://www.expressanalytics.com/blog/developing-machine-learning-models-for-dynamic-pricing/

FasterCapital. (023.). Time-Based Pricing: How to Charge Different Prices at Different Times. Retrieved from https://fastercapital.com/content/Time-Based-Pricing--How-to-Charge-Different-Prices-at-Different-Times.html

Tesauro, G., Das, R., Chan, H., Kephart, J. O., Lefurgy, C., Levine, D. W., & Rawson, F. (2020). Bayesian regression and Bitcoin. Springer Nature. Retrieved [date you accessed the website], from https://www.nature.com/articles/s43586-020-00001-2

Vendavo. (2023). Price Segmentation: What it is and How to Use it. Retrieved from https://www.vendavo.com/all/price-segmentation/

Ashrafi, B. (n.d.). Dynamic Pricing Using Machine Learning. Medium. Retrieved from https://medium.com/@baabak/dynamic-pricing-using-machine-learning-5e882282effe

AIMultiple. (n.d.). Dynamic Pricing Algorithm: A Comprehensive Guide. Retrieved [date you accessed the website], from https://research.aimultiple.com/dynamic-pricing-algorithm/

Ferreira, K. J., & Lee, B. (2015). The use of Machine Learning Algorithms in Recommender Systems: A Systematic Review. Expert Systems with Applications, 97, 205-227.

Bartodziej, C. J. (2017). The Concept Industry 4.0. In The Autonomy of the Person and the Digital Revolution (pp. 73-84). Springer Vieweg.

Rücker, A. Wes McKinney. (2011). Pandas: a foundational Python library for data analysis and statistics. Python for High Performance and Scientific Computing, 1-9.

Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... & Hassabis, D. (2015). Human-level control through deep reinforcement learning. Nature, 518(7540), 529-533.

Abadi, M., Barham, P., Chen, J., Chen, Z., Davis, A., Dean, J., ... & Ghemawat, S. (2016). TensorFlow: A System for Large-Scale Machine Learning. In OSDI (Vol. 16, pp. 265-283)

Chiang, W. C., Chen, J. C., & Xu, X. (2007). An overview of research on revenue management: current issues and future research. International Journal of Revenue Management, 1(1), 97-128.

Lee, J., Kao, H. A., & Yang, S. (2014). Service innovation and smart analytics for Industry 4.0 and big data environment. Procedia Cirp, 16, 3-8.

Chen, H., Chiang, R. H. L., & Storey, V. C. (2012). Business Intelligence and Analytics: From Big Data to Big Impact. MIS Quarterly, 36(4), 1165-1188.

Rossmann, B., Canzaniello, A., von der Gracht, H., & Hartmann, E. (2018). The future and social impact of Big Data Analytics in Supply Chain Management: Results from a Delphi study. Technological Forecasting and Social Change, 130, 135-149.

Zhang, D., Fjell, K., & Dröge, C. (2011). Optimal performance with real-time decision aids. Journal of Management Information Systems, 28(1), 251-286.

Mortensen, J. N., Tjøstheim, M., Løkken, I. K., Fossen, A. E., & Gondal, F. (2020). Policy Gradient Reinforcement Learning for Automated Blood Glucose Control in Type 1 Diabetes. Applied Sciences, 10(18), 6350. 1

Samuel, A. L. (1959). Some Studies in Machine Learning Using the Game of Checkers. IBM Journal of Research and Development, 3(3), 210-229.

Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., & Meger, D. (2018). Deep reinforcement learning that matters. In Proceedings of the AAAI Conference on Artificial Intelligence (Vol. 32, No. 1).

Lin, L. J. (1992). Self-improving reactive agents based on reinforcement learning, planning and teaching. Machine learning, 8(3), 293-321.

Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... & Petersen, S. (2015). Human-level control through deep reinforcement learning. Nature, 518(7540), 529-533.

Lillicrap, T. P., Hunt, J. J., Pritzel, A., Heess, N., Erez, T., Tassa, Y., Silver, D., & Wierstra, D. (2016). Continuous control with deep reinforcement learning. ArXiv, abs/1509.02971.

Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., & Meger, D. (2018). Deep reinforcement learning that matters. In Proceedings of the AAAI Conference on Artificial Intelligence (Vol. 32, No. 1).

Toon, Chris. (2019). Retrieved from https://towardsdatascience.com/deep-deterministic-policy-gradients-explained-2d94655a9b7b